

# Java's Place in the AI Revolution



Zoran Sevarac, *Deep Netts Technologies, Java Champion*



Frank Greco, *NYJavaSIG, Java Champion*



Pratik Patel, *Friends of OpenJDK (Foojay.io), Java Champion*



# *Takeaways*

*(yep... before the Giveaways)*

*Learn AI or your job/career is at risk*

*Java is a GREAT production language for AI*



*Models are probabilistic*

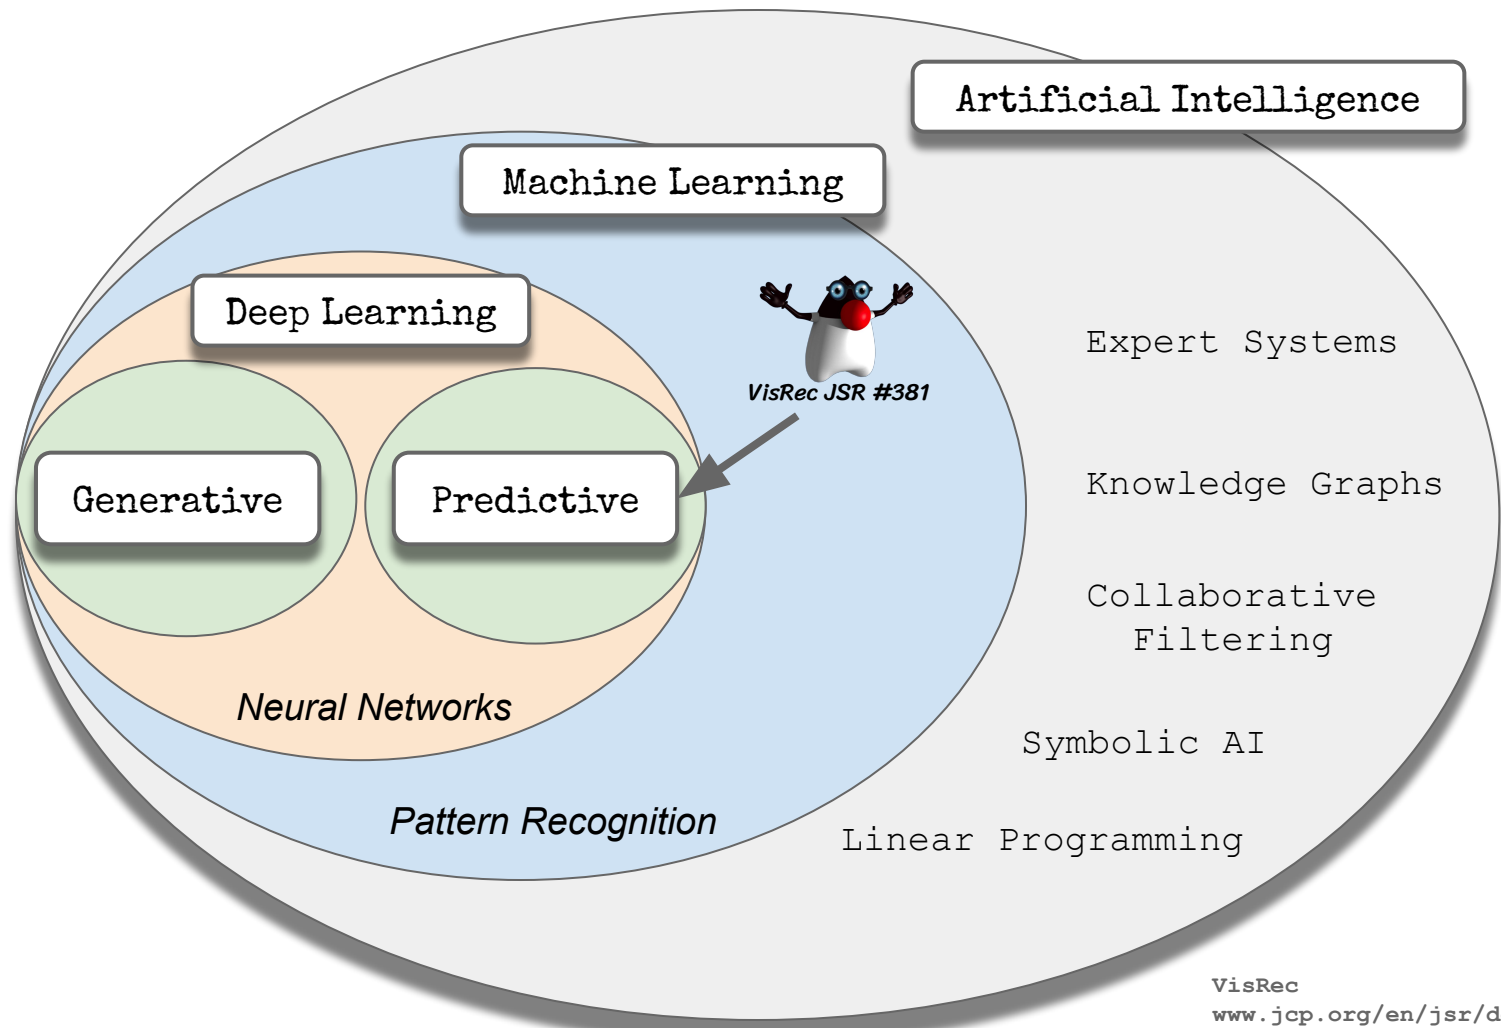
*You can use AI as a Java developer's tool  
or to develop AI Apps  
But learn how it works first...*



*“One of the most interesting aspects of the world is that it can be considered to be made up of patterns”*

Norbert Wiener (1948) - 1894-1964 - MIT





# *Predictive AI and Generative AI*

Most of AI \$\$ value (02/24) comes from PredAI: weather, image detection and classification, financial services, buying behavior, up/cross-selling...

PredAI has been deployed successfully for past 15+ yrs

PredAI is probably worth at least \$100B just to Google

GenAI typically used for more “creative”, content generation

GenAI growth and potential is huge. Market value may match PredAI in 3-5 yrs

# Value from AI technologies: Today → 3 years



Stanford

<https://www.youtube.com/watch?v=5p248yoa3oE>

July 26, 2023  
Andrew Ng

# ***Predictive AI (PredAI)***

forecasting  
*structured data*

Classification

Recommendation systems

Sales Forecasts

Fraud Detection

Predictive maintenance

# ***Generative AI (GenAI)***

creation  
*unstructured data*

Creative/Design

Ideation


Prototyping

Simulation

Translation / Summarization


# **Predictive AI (PredAI)**

forecasting  
structured data



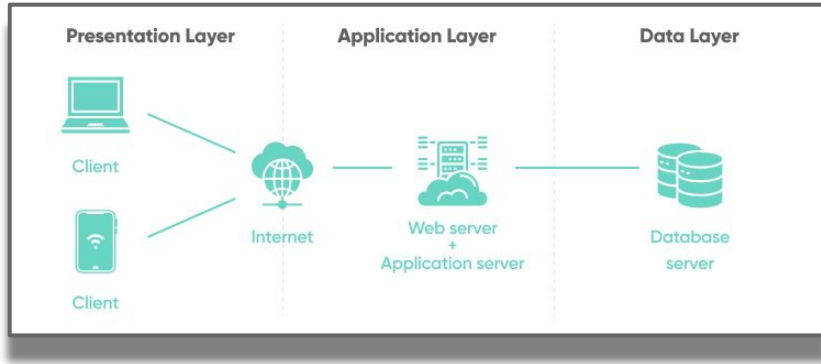
# **Generative AI (GenAI)**

creation  
unstructured data

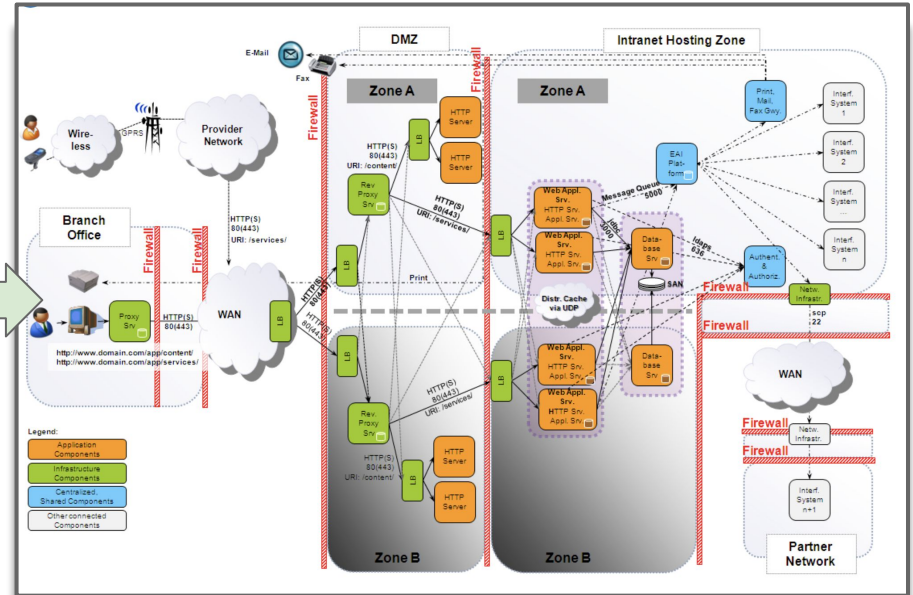

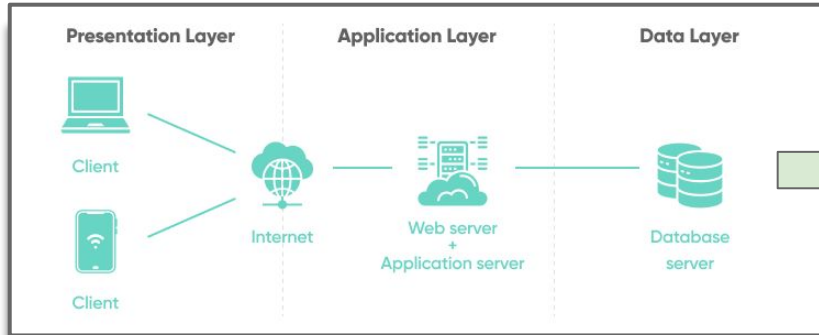





# ***Evolving Architectures***



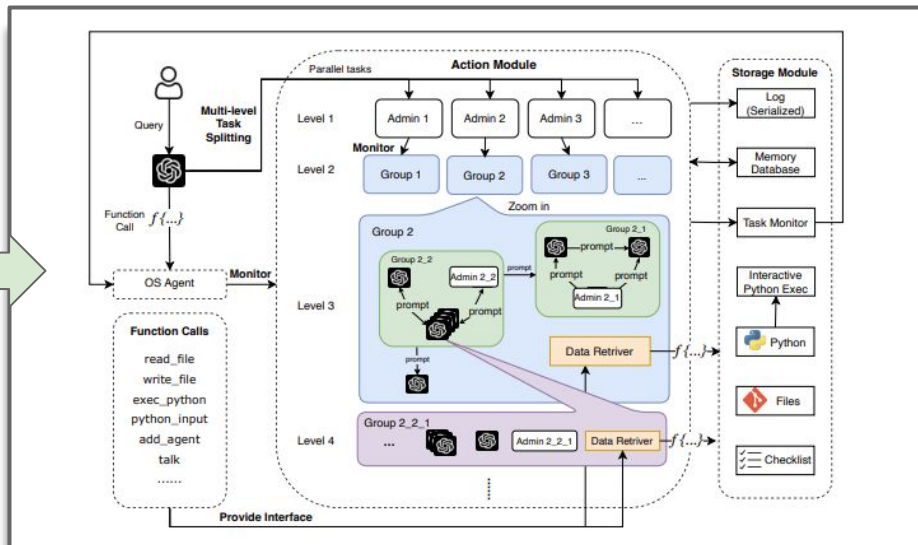

# Web



# Web



# GenAI



# GenAI

# Why Use Java for Production AI Applications?

ChatGPT 4o ▾

🔗 Share

## Top 10 Reasons to Use Java for Production AI Applications (Enterprise Focused)

1. **Seamless Integration** – Java fits natively into existing enterprise stacks (Spring, Kafka, Hadoop, etc.).
2. **Scalability** – Proven JVM performance for large-scale, multi-threaded applications.
3. **Robust Tooling** – Mature build, monitoring, and deployment tools (Maven, JFR, JMX).
4. **Security** – Strong, enterprise-grade security features baked into the platform.
5. **Maintainability** – Static typing and strong IDE support enhance long-term maintainability.
6. **Cross-Platform Deployment** – JVM portability ensures consistent behavior across environments.
7. **Enterprise Support** – Backed by long-term support from Oracle, Red Hat, and others.
8. **Developer Pool** – Large pool of experienced Java developers in enterprise settings.
9. **Memory Management** – Sophisticated garbage collection options tailored for performance.
10. **AI Ecosystem Growth** – Expanding support via libraries like **Tribuo**, **Deep Java Library (DJL)**, and the standardized **JSR 381 Visual Recognition API**.

langchain4j / langchain4j

Type to search

<> Code

Issues 220

Pull requests 64

Discussions

Actions

Projects

Security

Insights

langchain4j

Public

Watch 72

Fork 775

Star 4k

main

23 Branches

35 Tags

Go to file

t

Add file

<> Code

langchain4j

re #725: PR 5: Updated documentation. (#1473) ✓

fdb052 · 7 hours ago

1,014 Commits

.devcontainer	Add Dev Container support (#337)	7 months ago
.github	Update pull_request_template.md	2 weeks ago
.mvn/wrapper	Correctly configure Maven wrapper (#348)	7 months ago
code-execution-engines	changed version to 0.33.0-SNAPSHOT	last week
docker/ollama	feat : create llama3 model image (#1083)	2 months ago
docs	re #725: PR 5: Updated docum	
document-loaders	changed version to 0.33.0-SNA	
document-parsers	changed version to 0.33.0-SNAPSHOT	last week
embedding-store-filter-parsers/langchain4j-...	changed version to 0.33.0-SNAPSHOT	last week
experimental/langchain4j-experimental-sql	changed version to 0.33.0-SNAPSHOT	last week
langchain4j-anthropic	changed version to 0.33.0-SNAPSHOT	last week
langchain4j-azure-ai-search	Feat(#1383): mutualise EmbeddingMatches handling (#1...	last week
langchain4j-azure-cosmos-mongo-vcore	changed version to 0.33.0-SNAPSHOT	last week
langchain4j-azure-cosmos-nosql	changed version to 0.33.0-SNAPSHOT	last week
langchain4j-azure-open-ai	changed version to 0.33.0-SNAPSHOT	last week

About

Java version of LangChain

[docs.langchain4j.dev](#)

java

embeddings

gemini

openai

chroma

llama

gpt

pinecone

onnx

weaviate

huggingface

milvus

vector-database

openai-api

chatgpt

langchain

anthropic

pgvector

ollama

Activity

4k stars

72 watching

775 forks

Report repository

Releases 34

0.32.0 Latest

2 weeks ago

+ 33 releases

deepnetts / deepnetts-communityedition

Code Pull requests 1 Actions Projects Security Insights

deepnetts-communityedition Public

forked from sevarac/deepnetts-communityedition

Watch 3 Fork 9 Star 49

community-visrec 1 Branch

This branch is 21 commits ahead of sevarac/deepnetts-communityedition

neuroph iris flowers example comments and normalization added f644646 · 3 years ago 235 Commits

View license Activity 49 stars 3 watching 9 forks

.github/workflows Temporarily disabling RandomLinearDataGenerator due t... 4 years ago

deepnetts-core fixed serialVersionUID issue using transient inputTensor 3 years ago

deepnetts-examples iris flowers example comments

logs Removing getTargetNames() from 4 years ago

notebooks Moving linear.csv to the dataset

.gitignore Without this changes the net file

COPYING Update COPYING

README.md Merge pull request #12 from ge

nbactions.xml Added regularization

pom.xml Removing getTargetNames() from

README License

<https://github.com/deepnetts/deepnetts-communityedition>

DEEP NETTS

HOME PRODUCT DOCUMENTATION BLOG COMPANY LICENSING

FREE DOWNLOAD

# Empowering AI Revolution

AI Development Platform in Java

Deep Netts platform accelerates AI adoption by simplifying development, integration and deployment, making predictive AI more accessible and efficient for Java-based enterprises.

## Visual AI Builder

## Deep Learning Java Library

```
// create an instance of a neural network using builder
FeedForwardNetwork neuralNet = FeedForwardNetwork.builder()
    .addInputLayer(4)
    .addFullyConnectedLayer(32, ActivationType.TANH)
    .addOutputLayer(3, ActivationType.SOFTMAX)
    .lossFunction(LossType.CROSS_ENTROPY)
    .randomSeed(123)
    .build();
```



# How to use AI models in Java

## 1. Use Web service - serve model as web service

You can build a model in any language, using a model is language agnostic, potential issues with latency and scalability.

## 2. Use wrappers or native libraries from Java

Native dependencies, limited scalability, memory issues, distribution and maintenance overhead (aka. nightmare)

## 3. Use Java native AI libraries

Highly scalable, low latency, easy to use, integrate with existing development and production environment, and distribute on large scale.

Out of the box models mostly not available.



# Overview of Java AI Libraries

	Description	Pros	Cons
<b>Tensorflow Java API</b>	Java API for Tensorflow developed by community	Many out of the box models available, GPU and TPU support, many algorithms and architectures,	Large size, Not covered with compatibility guarantees Requires lower level understanding of TF internals
<b>DJL</b>	Wrapper around Python based frameworks	Many out of the box models available	Mixed tech stacks
<b>DL4J</b>	Java Deep Learning API on top of native numeric libraries	High performance, GPU support many algorithms and architectures, feature rich	Native dependencies, Large size, not actively developed, complex to learn and use
<b>Deep Netts</b>	Java native, gpu support through jcuda	Easy to learn and use, Easy to integrate into existing Java infrastructure , highly scalable, low latency	Not all algorithms are supported

# About Deep Netts AI Platform for Java

## Deep Netts provides:

- **Deep Learning Java Library**  
Build, train, and deploy NN natively in Java
- **Visual AI Builder Tool**  
Intuitive, no-code/low-code environment for model creation
- **Professional Support**  
Expert assistance for development and production needs
- **Community edition**  
Free and open access for exploration and learning

## Deep Netts makes predictive AI

- ✓ More accessible
- ✓ More efficient
- ✓ Easier to integrate and deploy within Java-based enterprise systems

## By accelerating and simplifying:

- AI development
- AI integration
- AI deployment


**Vision:** We would like to see the Java platform evolve to meet modern AI requirements.



# AI Demo using Deep Netts

1. Define the problem description
2. Prepare the data
3. Build the model
4. Test/evaluate the model
5. Use the model


# 1. Define the problem description: Fraud Detection



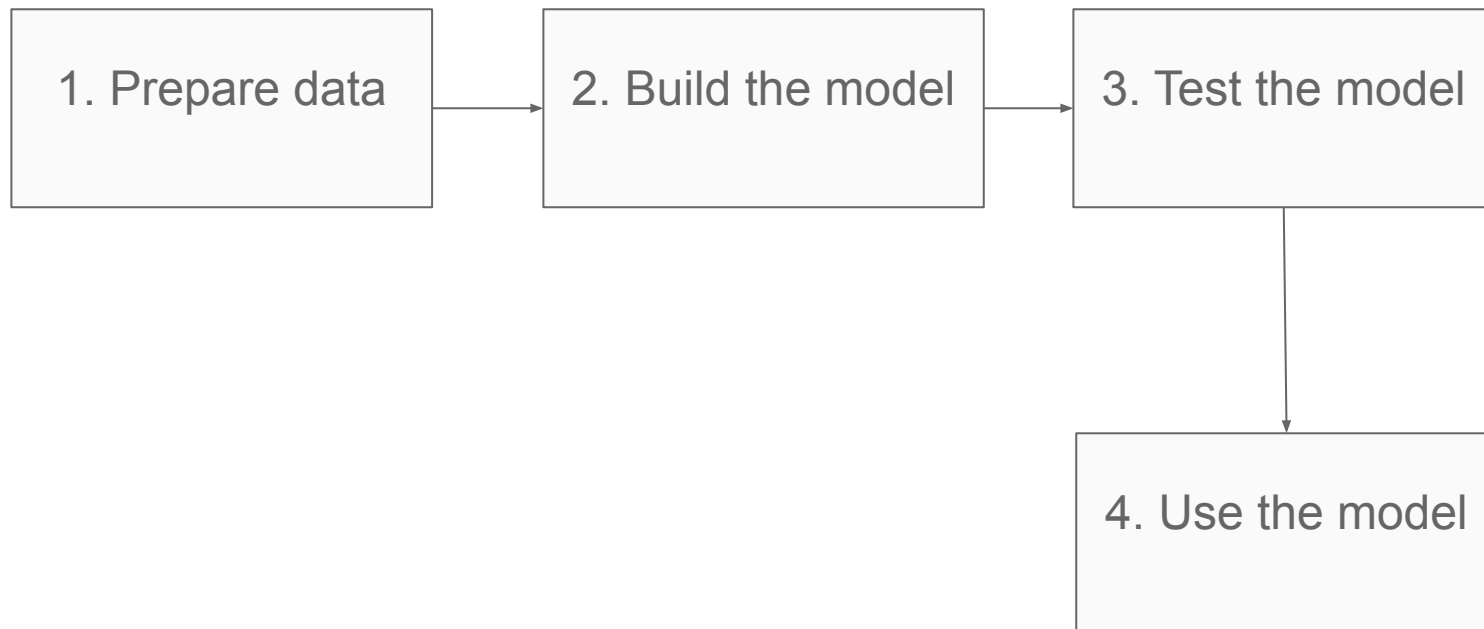
Given the credit card transaction details, how likely is that it is a fraud?

# AI/ML Models - Software Engineer definition

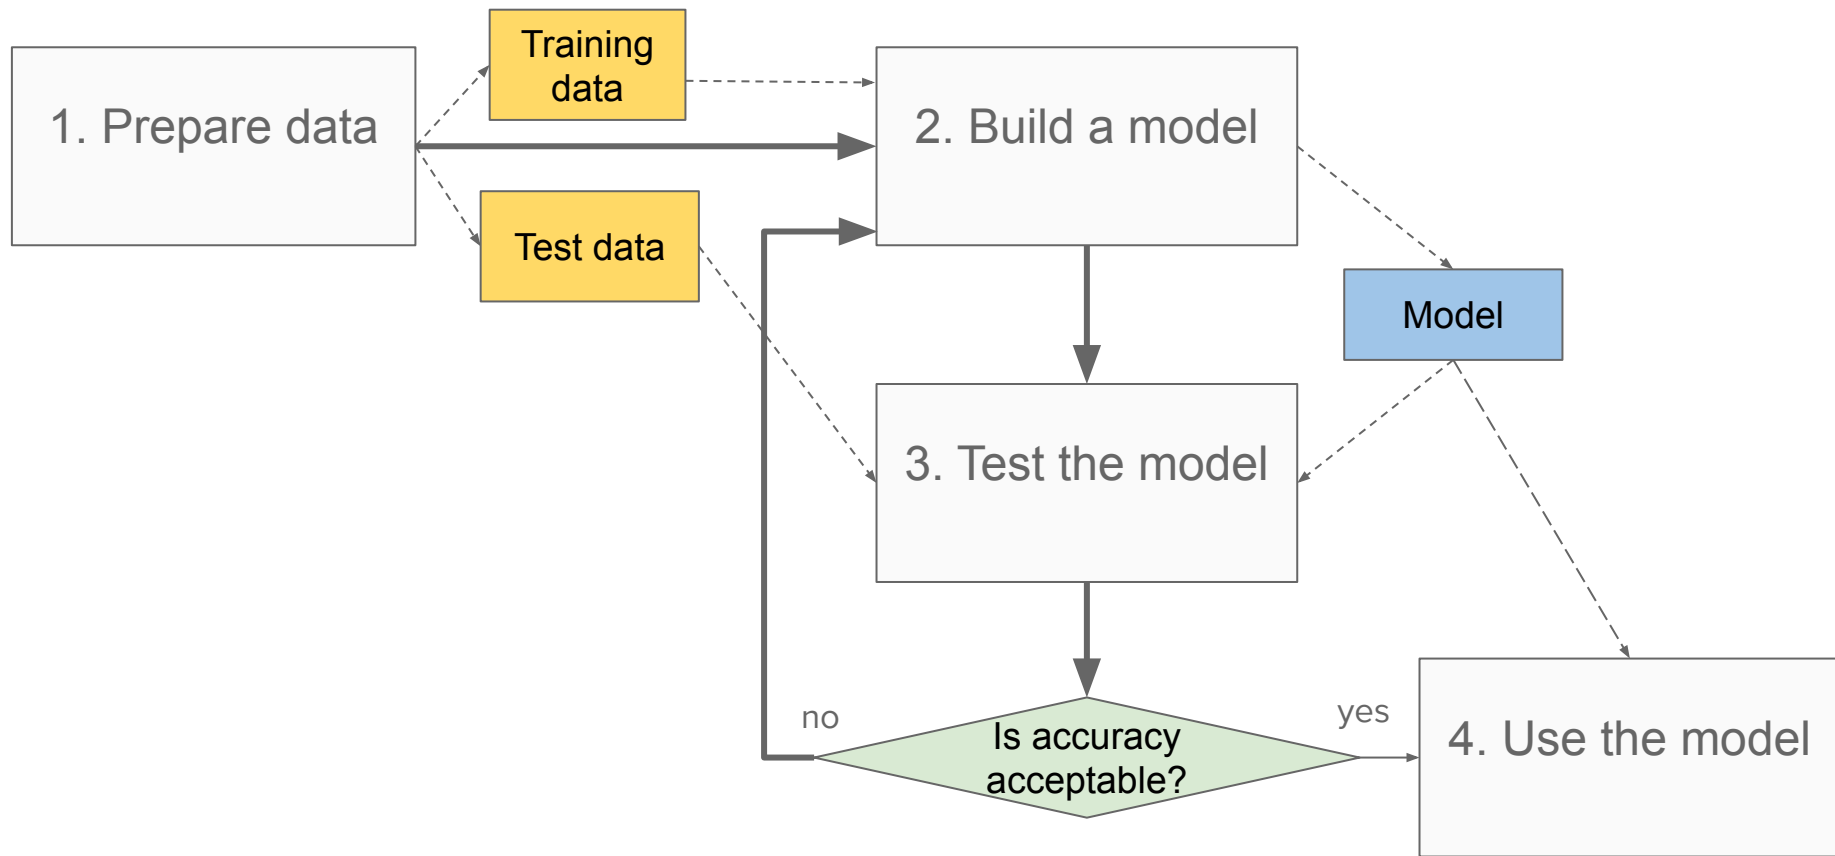
Model is just another type of abstraction in developer's toolbox.



# Workflow Overview



# Workflow Overview





## 2. Prepare the data

### // specify CSV file options

```
CsvReadOptions.Builder builder =  
CsvReadOptions.builder("creditcard.csv")  
    .separator(',') // values are coma-delimited  
    .header(true); // first line contains column names  
CsvReadOptions options = builder.build();
```

Load data from CSV file into a  
**Data Frame** - basically a Table  
using Tablesaw Java Library

### // load data into a data frame

```
Table dataTable = Table.read().usingOptions(options);
```

### // prepare data for training

```
DataPreparation dataPrep = new DataPreparation(dataTable);
```

Out of the box, single method  
call, data utilities

### // print header and first few rows to see what's loaded

```
dataPrep.previewRows(5);
```

Check loaded - rows and  
columns

### // print columns with corresponding types

```
dataPrep.columnInfo();
```

# Preprocess the data to ensure data quality

```
dataTable.removeColumns("Time");  
  
dataTable.dropDuplicateRows();  
  
// check if there are any missing values  
dataPrep.countMissingValues();  
  
dataPrep.handleMissingValues();  
  
// examine basic statistics summary  
dataPrep.statistics();  
  
// check if data set is balanced  
DataPreparation.checkClassBalance(dataTable);  
  
// create balanced subset  
Table balancedData =  
DataPreparation.createBalancedSample(dataTable, "Class", 1)
```

Decide which columns are needed

Remove duplicate rows



Check and fix missing values

Explore statistics for each column, check distribution

Check for outliers

Make sure data set is balanced

# Balancing Data Set



# Prepare data for training

```
// create data set for neural network training
```

```
TabularDataSet dataSet = DataPreparation.createDataSet(balancedData);
```

Create data set for  
nn training

```
// scale data to value range [0, 1]
```

```
DataSets.scaleToMax(dataSet);
```

Scale values to  
range used in nn

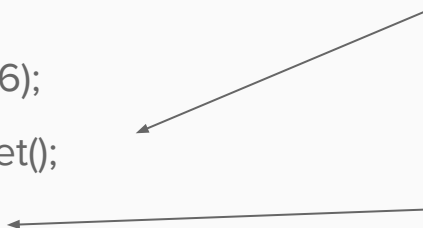
```
// split data into training and test set
```

```
TrainTestSplit split = DataSets.trainTestSplit(dataSet, 0.6);
```

```
DataSet<MLDataItem> trainingSet = split.getTrainingSet();
```

```
DataSet<MLDataItem> testSet = split.getTestSet();
```

Training set to train  
the model



Test set to estimate  
accuracy on  
unseen data

### 3. Build a Model - Feed Forward Neural Network

**// instantiate and configure a neuralnet for binary classification**


```
FeedForwardNetwork neuralNet = FeedForwardNetwork.builder()
    .addInputLayer(numInputs)
    .addFullyConnectedLayer(32, ActivationType.TANH)
    .addOutputLayer(numOutputs, ActivationType.SIGMOID)
    .lossFunction(LossType.CROSS_ENTROPY)
    .build();
```

**// set parameters of the training algorithm**

```
neuralNet.getTrainer().setStopError(0.02f)
                    .setStopEpochs(10000)
                    .setLearningRate(0.001f);
```

**// train the model prepared data set**

```
neuralNet.train(trainingSet);
```



# Configure the Model Architecture

```
FeedForwardNetwork neuralNet =
```

```
FeedForwardNetwork.builder()
```

```
.addInputLayer(29)
```

```
.addFullyConnectedLayer(32, ActivationType.TAN
```

```
.addOutputLayer(1, ActivationType.SIGMOID)
```

```
.lossFunction(LossType.CROSS_ENTROPY)
```

```
.build();
```

How many inputs

Hidden layer units  
and activation

Single output, for  
binary classification

Loss function for  
binary classification

# Training Configuration

```
// set parameters of the training algorithm  
neuralNet.getTrainer()
```

```
.setStopError(0.02f)
```

```
.setStopEpochs(10000)
```

```
.setLearningRate(0.01f)
```

```
// start training with specified data set  
neuralNet.train(trainingSet);
```

Get instance of training algo

Training stops when this error is reached

Or this number of iterations/epochs is reached

Step size for adjusting weights  
~1% of error

# Training Process

```
while(epoch < stopEpochs && error > stopError) {  
    for each example in trainingSet {  
        calculate prediction for input  
        calculate prediction error  
        tune weights in order to reduce error  
    }  
}
```



# Debugging and Hyperparameter Tuning

Q: What if error is increasing during the training?

Try smaller learning rate

Q: What if error is not decreasing even after many training epochs?

Try adding more hidden units and layers

Q: Is the model overfitting (low training, high test error)?

Add more data or try simpler architecture

Hint: Use Visual AI Builder tool to run experiments and tune settings

## 4. Testing/Evaluating the model - How accurate?

```
EvaluationMetrics em = neuralNet.test(testSet);
```

Calculate various  
eval metrics using  
test set


**Metrics:** True Positive, True Negative, False positive, False Negative  
Accuracy, Precision, Recall, F-Score

1. How often the prediction is correct when the model predicts fraud?
2. How often the prediction is incorrect when the model predicts fraud?  
How often the prediction is correct in total? (accuracy)
3. What is more important/expensive (False Positive or False Negative)?  
The trick is to find balance.

## 5. Using the model - fraud probability


What is that the probability that the given transaction is fraud?

```
Tensor prediction = neuralNet.predict(inputTensor);  
if (prediction.get(0) > 0.8) {  
    System.out.println("There is a high probability that this is fraud!");  
}
```



Feed model with input as tensor and get prediction as tensor

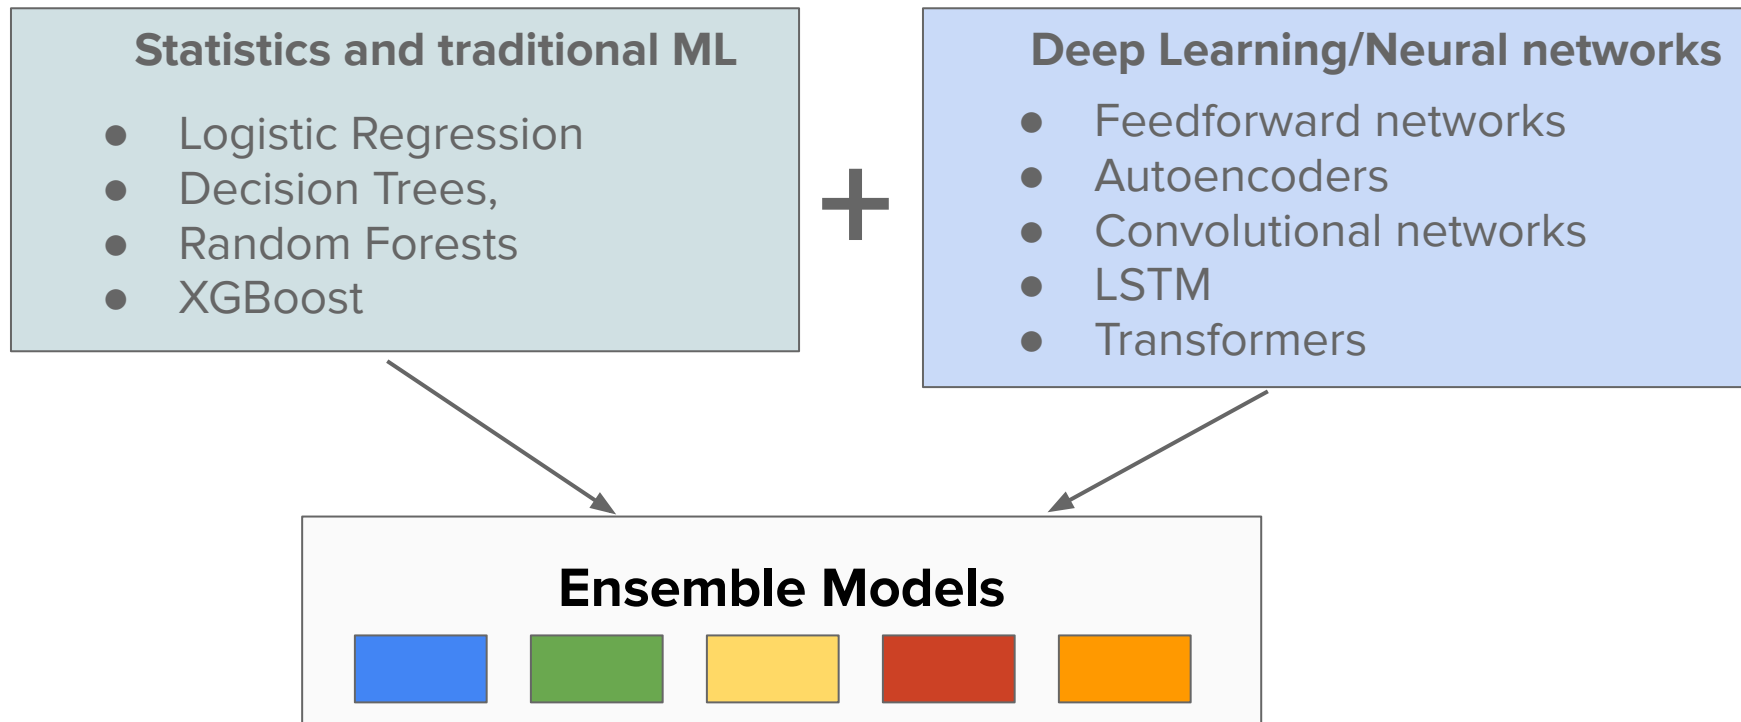
### Using JSR381



Use model through higher level classifier API

```
BinaryClassifier<float[]> binClassifier = new FeedForwardNetBinaryClassifier(neuralNet);  
Float probability = binClassifier.classify(testTransaction);
```

# Combine with Other Methods



# Take Your Java AI Projects to the Next Level with Deep Netts Expertise!

- ◆ **Are you using AI in your Java applications?**

- What libraries or tools are you currently using?
- Are you facing challenges building or deploying your AI models
- Have an exciting project you'd like to discuss?

- ◆ **Deep Netts offers expert consulting for:**

- AI/ML solution design and integration in Java
- Optimizing AI models for production environments
- Scaling AI across distributed Java systems

# What's next for Java AI developers

## Resources:

- <https://github.com/deepnetts/CreditCardFraudDetection> Demo on Github
- <https://www.deepnetts.com/download-latest> Free Deep Netts Download -Free for development, prototyping and education

## Join Linkedin Group AI/ML for Java Ecosystem

<https://www.linkedin.com/groups/10084933/>

## Quick Survey for future webinars

<https://forms.gle/RfskNfohaKkhP15Q9>